Recommendations_with_IBM

March 1, 2021

1 Recommendations with IBM

In this notebook, you will be putting your recommendation skills to use on real data from the IBM Watson Studio platform.

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project RUBRIC. **Please save regularly.**

By following the table of contents, you will build out a number of different methods for making recommendations that can be used for different situations.

1.1 Table of Contents

I. Section **??** II. Section **??** III. Section **??** IV. Section **??** V. Section **??** VI. Section **??**

At the end of the notebook, you will find directions for how to submit your work. Let's get started by importing the necessary libraries and reading in the data.

```
In [1]: import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import project_tests as t
        import pickle
        import seaborn as sns
        from collections import defaultdict

        %matplotlib inline

        df = pd.read_csv('data/user-item-interactions.csv')
        df_content = pd.read_csv('data/articles_community.csv')
        del df['Unnamed: 0']
        del df_content['Unnamed: 0']

        # Show df to get an idea of the data
        df.head()
```

```
Out[1]:    article_id                                              title  \
        0      1430.0  using pixiedust for fast, flexible, and easier...
        1      1314.0       healthcare python streaming application demo
        2      1429.0         use deep learning for image classification
```

```
3      1338.0          ml optimization using cognitive assistant
4      1276.0          deploy your python model as a restful api

                                      email
0  ef5f11f77ba020cd36e1105a00ab868bbdbf7fe7
1  083cbdfa93c8444beaa4c5f5e0f5f9198e4f9e0b
2  b96a4f2e92d8572034b1e9b28f9ac673765cd074
3  06485706b34a5c9bf2a0ecdac41daf7e7654ceb7
4  f01220c46fc92c6e6b161b1849de11faacd7ccb2
```

In [2]:
```python
# Show df_content to get an idea of the data
df_content.head()
```

Out[2]:
```
                                            doc_body  \
0  Skip navigation Sign in SearchLoading...\r\n\r...
1  No Free Hunch Navigation * kaggle.com\r\n\r\n ...
2   * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
3  DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...
4  Skip navigation Sign in SearchLoading...\r\n\r...

                                     doc_description  \
0  Detect bad readings in real time using Python ...
1  See the forest, see the trees. Here lies the c...
2  Heres this weeks news in Data Science and Bi...
3  Learn how distributed DBs solve the problem of...
4  This video demonstrates the power of IBM DataS...

                                       doc_full_name doc_status  article_id
0  Detect Malfunctioning IoT Sensors with Streami...       Live           0
1  Communicating data science: A guide to present...       Live           1
2           This Week in Data Science (April 18, 2017)       Live           2
3  DataLayer Conference: Boost the performance of...       Live           3
4      Analyze NY Restaurant data using Spark in DSX       Live           4
```

1.1.1 Part I : Exploratory Data Analysis

Use the dictionary and cells below to provide some insight into the descriptive statistics of the data.

1. What is the distribution of how many articles a user interacts with in the dataset? Provide a visual and descriptive statistics to assist with giving a look at the number of times each user interacts with an article.

In [3]:
```python
# Fill in the median and maximum number of user_article interactios below

# 50% of individuals interact with ____ number of articles or fewer.
median_val = df.groupby('email')['article_id'].count().median()
# The maximum number of user-article interactions by any 1 user is ______.
max_views_by_user = df.groupby('email')['article_id'].count().max()
```

```
In [4]: # Distribution plot of number of user-article interactios

        sns.set_style('whitegrid')
        sns.displot(df.groupby('email')['article_id'].count(), kind='kde')
        plt.xlabel('Number of user-article interactios')
        plt.title('Distribution of of how many articles a user interacts with.')
```

```
Out[4]: Text(0.5, 1.0, 'Distribution of of how many articles a user interacts with.')
```

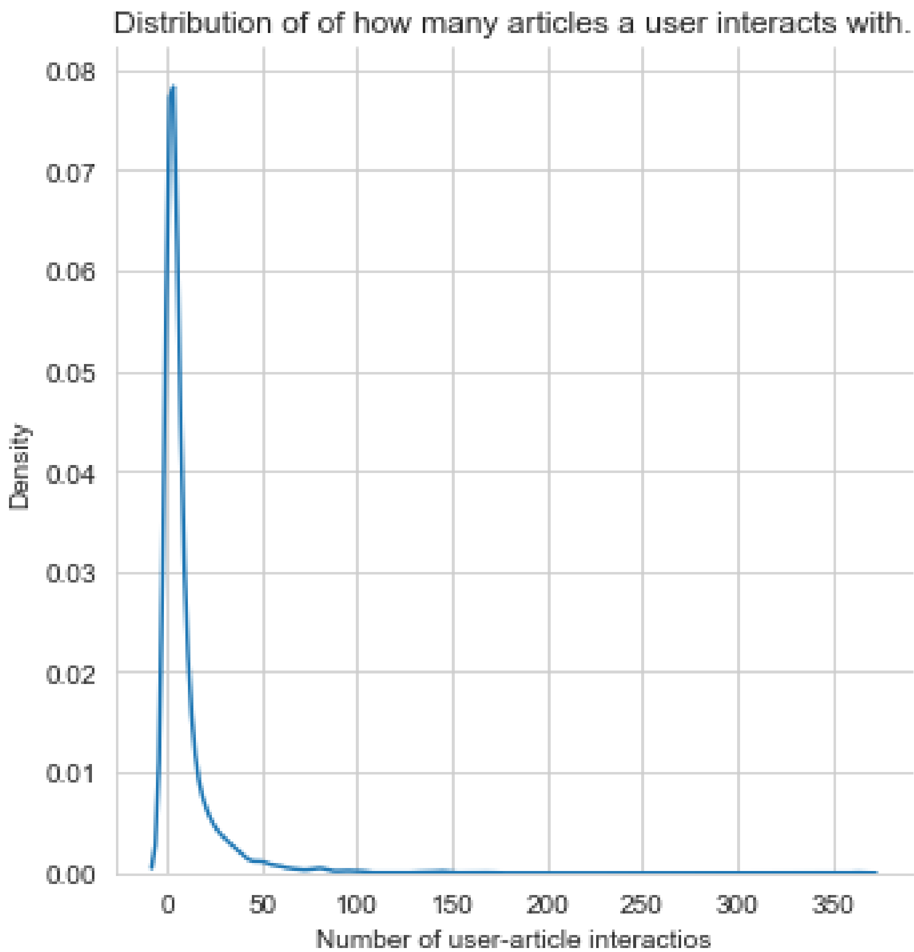

2. Explore and remove duplicate articles from the **df_content** dataframe.

```
In [5]: # shape of df_content dataframe
        df_content.shape
```

```
Out[5]: (1056, 5)
```

```
In [6]: # number of duplicate rows
        df_content.duplicated('article_id').sum()
```

Out[6]: 5

In [7]:
```python
# Find and explore duplicate articles
df_content[df_content.duplicated('article_id', keep=False)]
```

Out[7]:
```
                                              doc_body  \
50   Follow Sign in / Sign up Home About Insight Da...
221  * United States\r\n\r\nIBMő * Site map\r\n\r\n...
232  Homepage Follow Sign in Get started Homepage *...
365  Follow Sign in / Sign up Home About Insight Da...
399  Homepage Follow Sign in Get started * Home\r\n...
578  This video shows you how to construct queries ...
692  Homepage Follow Sign in / Sign up Homepage * H...
761  Homepage Follow Sign in Get started Homepage *...
970  This video shows you how to construct queries ...
971  Homepage Follow Sign in Get started * Home\r\n...

                                       doc_description  \
50                           Community Detection at Scale
221  When used to make sense of huge amounts of con...
232  If you are like most data scientists, you are ...
365  During the seven-week Insight Data Engineering...
399  Todays world of data science leverages data f...
578  This video shows you how to construct queries ...
692  One of the earliest documented catalogs was co...
761  Todays world of data science leverages data f...
970  This video shows you how to construct queries ...
971  If you are like most data scientists, you are ...

                                         doc_full_name doc_status  article_id
50                         Graph-based machine learning       Live          50
221  How smart catalogs can turn the big data flood...       Live         221
232  Self-service data preparation with IBM Data Re...       Live         232
365                        Graph-based machine learning       Live          50
399  Using Apache Spark as a parallel processing fr...       Live         398
578                               Use the Primary Index       Live         577
692  How smart catalogs can turn the big data flood...       Live         221
761  Using Apache Spark as a parallel processing fr...       Live         398
970                               Use the Primary Index       Live         577
971  Self-service data preparation with IBM Data Re...       Live         232
```

In [8]:
```python
# Remove any rows that have the same article_id - only keep the first
df_content.drop_duplicates('article_id', inplace=True)
```

In [9]:
```python
# Check shape of df_content after dropping duplicates
df_content.shape
```

Out[9]: (1051, 5)

3. Use the cells below to find:
a. The number of unique articles that have an interaction with a user.
b. The number of unique articles in the dataset (whether they have any interactions or not). **c.** The number of unique users in the dataset. (excluding null values) **d.** The number of user-article interactions in the dataset.

```
In [10]: # Number of unique articles that have an interaction with a user
         df.groupby('article_id')['email'].count().sort_values()

Out[10]: article_id
         1092.0      1
         984.0       1
         417.0       1
         1237.0      1
         1233.0      1
                  ...
         1364.0    627
         1427.0    643
         1431.0    671
         1330.0    927
         1429.0    937
         Name: email, Length: 714, dtype: int64

In [11]: # The number of unique articles on the IBM platform
         df_content['article_id'].nunique()

Out[11]: 1051

In [12]: # The number of unique users
         df['email'].nunique()

Out[12]: 5148

In [13]: # The number of user-article interactions
         df.shape

Out[13]: (45993, 3)

In [14]: # The number of user-article interactions after dropping duplicates
         df.drop_duplicates().shape[0]

Out[14]: 33682

In [15]: unique_articles = 714  # The number of unique articles that have at least one interacti
         total_articles = 1051  # The number of unique articles on the IBM platform
         unique_users = 5148  # The number of unique users
         user_article_interactions = 45993  # The number of user-article interactions
```

4. Use the cells below to find the most viewed **article_id**, as well as how often it was viewed. After talking to the company leaders, the `email_mapper` function was deemed a reasonable way to map users to ids. There were a small number of null values, and it was found that all of these null values likely belonged to a single user (which is how they are stored using the function below).

In [16]:
```
df.groupby('article_id')['email'].count().sort_values(ascending=False)
```

Out[16]:
```
article_id
1429.0    937
1330.0    927
1431.0    671
1427.0    643
1364.0    627
         ...
1113.0      1
1119.0      1
984.0       1
1127.0      1
1266.0      1
Name: email, Length: 714, dtype: int64
```

In [17]:
```
df.groupby('article_id')['email'].count().idxmax()
```

Out[17]:
```
1429.0
```

In [18]:
```
# The most viewed article in the dataset as a string with one value following the decim
most_viewed_article_id = "1429.0"
# The most viewed article in the dataset was viewed how many times?
max_views = df.groupby('article_id')['email'].count().max()
```

In [19]:
```
# No need to change the code here - this will be helpful for later parts of the noteboo
# Run this cell to map the user email to a user_id column and remove the email column

def email_mapper():
    coded_dict = dict()
    cter = 1
    email_encoded = []

    for val in df['email']:
        if val not in coded_dict:
            coded_dict[val] = cter
            cter += 1

        email_encoded.append(coded_dict[val])
    return email_encoded


email_encoded = email_mapper()
del df['email']
df['user_id'] = email_encoded

# show header
df.head()
```

```
Out[19]:    article_id                                              title  user_id
         0      1430.0  using pixiedust for fast, flexible, and easier...        1
         1      1314.0       healthcare python streaming application demo        2
         2      1429.0         use deep learning for image classification        3
         3      1338.0          ml optimization using cognitive assistant        4
         4      1276.0          deploy your python model as a restful api        5
```

```
In [20]: # If you stored all your results in the variable names above,
         # you shouldn't need to change anything in this cell

         sol_1_dict = {
             '`50% of individuals have _____ or fewer interactions.`': median_val,
             '`The total number of user-article interactions in the dataset is ______.`': user_a
             '`The maximum number of user-article interactions by any 1 user is ______.`': max_v
             '`The most viewed article in the dataset was viewed _____ times.`': max_views,
             '`The article_id of the most viewed article is ______.`': most_viewed_article_id,
             '`The number of unique articles that have at least 1 rating ______.`': unique_artic
             '`The number of unique users in the dataset is ______`': unique_users,
             '`The number of unique articles on the IBM platform`': total_articles
         }

         # Test your dictionary against the solution
         t.sol_1_test(sol_1_dict)
```

```
It looks like you have everything right here! Nice job!
```

1.1.2 Part II: Rank-Based Recommendations

Unlike in the earlier lessons, we don't actually have ratings for whether a user liked an article or not. We only know that a user has interacted with an article. In these cases, the popularity of an article can really only be based on how often an article was interacted with.

1. Fill in the function below to return the **n** top articles ordered with most interactions as the top. Test your function using the tests below.

```
In [21]: df.head()
```

```
Out[21]:    article_id                                              title  user_id
         0      1430.0  using pixiedust for fast, flexible, and easier...        1
         1      1314.0       healthcare python streaming application demo        2
         2      1429.0         use deep learning for image classification        3
         3      1338.0          ml optimization using cognitive assistant        4
         4      1276.0          deploy your python model as a restful api        5
```

```
In [22]: def get_top_articles(n, df=df):
             '''
             INPUT:
             n - (int) the number of top articles to return
             df - (pandas dataframe) df as defined at the top of the notebook
```

```
        OUTPUT:
        top_articles - (list) A list of the top 'n' article titles

        '''

        top_article_ids = get_top_article_ids(n)
        top_articles = list(df.groupby('article_id')[
                                   'title'].unique().apply(lambda t: t[0])[top_article_ids])
        # alternately
        # top_articles = list(df[df['article_id'].isin(top_article_ids)].drop_duplicates('a

        # Return the top article titles from df (not df_content)
        return top_articles


    def get_top_article_ids(n, df=df):
        '''
        INPUT:
        n - (int) the number of top articles to return
        df - (pandas dataframe) df as defined at the top of the notebook

        OUTPUT:
        top_article_ids - (list) A list of the top 'n' article ids

        '''

        top_article_ids = list(df.groupby('article_id')[
                                  'user_id'].count().sort_values(ascending=False).index[:n])
        # alternative method
        # top_article_ids = list(df['article_id'].value_counts().index[:5])

        return top_article_ids  # Return the top article ids
```

In [23]:
```
print(get_top_articles(10))
print(get_top_article_ids(10))
```

```
['use deep learning for image classification', 'insights from new york car accident reports', 'v
[1429.0, 1330.0, 1431.0, 1427.0, 1364.0, 1314.0, 1293.0, 1170.0, 1162.0, 1304.0]
```

In [24]:
```
# Test your function by returning the top 5, 10, and 20 articles
top_5 = get_top_articles(5)
top_10 = get_top_articles(10)
top_20 = get_top_articles(20)

# Test each of your three lists from above
t.sol_2_test(get_top_articles)
```

```
Your top_5 looks like the solution list! Nice job.
Your top_10 looks like the solution list! Nice job.
Your top_20 looks like the solution list! Nice job.
```

1.1.3 Part III: User-User Based Collaborative Filtering

1. Use the function below to reformat the **df** dataframe to be shaped with users as the rows and articles as the columns.

- Each **user** should only appear in each **row** once.
- Each **article** should only show up in one **column**.
- **If a user has interacted with an article, then place a 1 where the user-row meets for that article-column**. It does not matter how many times a user has interacted with the article, all entries where a user has interacted with an article should be a 1.
- **If a user has not interacted with an item, then place a zero where the user-row meets for that article-column**.

Use the tests to make sure the basic structure of your matrix matches what is expected by the solution.

```
In [25]: # create the user-article matrix with 1's and 0's

         def create_user_item_matrix(df):
             '''
             INPUT:
             df - pandas dataframe with article_id, title, user_id columns

             OUTPUT:
             user_item - user item matrix

             Description:
             Return a matrix with user ids as rows and article ids on the columns with 1 values
             an article and a 0 otherwise
             '''

             user_item = df.groupby(['user_id', 'article_id']
                                    ).count().unstack().replace(np.nan, 0)
             user_item[user_item > 0] = 1
             user_item.columns = user_item.columns.droplevel()
             # alternatively
             # user_item.columns = user_item.columns.get_level_values(1)

             return user_item  # return the user_item matrix


         user_item = create_user_item_matrix(df)
```

```
In [26]: # Tests: You should just need to run this cell.  Don't change the code.
         assert user_item.shape[0] == 5149, "Oops!  The number of users in the user-article matr
         assert user_item.shape[1] == 714, "Oops!  The number of articles in the user-article ma
         assert user_item.sum(axis=1)[
             1] == 36, "Oops!  The number of articles seen by user 1 doesn't look right."
         print("You have passed our quick tests!  Please proceed!")
```

```
You have passed our quick tests!  Please proceed!
```

2. Complete the function below which should take a user_id and provide an ordered list of the most similar users to that user (from most similar to least similar). The returned result should not contain the provided user_id, as we know that each user is similar to him/herself. Because the results for each user here are binary, it (perhaps) makes sense to compute similarity as the dot product of two users.

Use the tests to test your function.

```
In [27]: def find_similar_users(user_id, user_item=user_item):
             '''
             INPUT:
             user_id - (int) a user_id
             user_item - (pandas dataframe) matrix of users by articles:
                         1's when a user has interacted with an article, 0 otherwise

             OUTPUT:
             similar_users - (list) an ordered list where the closest users (largest dot product
                             are listed first

             Description:
             Computes the similarity of every pair of users based on the dot product
             Returns an ordered

             '''

             # compute similarity of each user to the provided user
             user_similarity = user_item.loc[user_id].dot(user_item.transpose())

             # sort by similarity
             user_similarity_sorted = user_similarity.sort_values(ascending=False)

             # create list of just the ids
             most_similar_users = user_similarity_sorted.index

             # remove the own user's id
             most_similar_users = most_similar_users.drop(user_id)

             # return a list of the users in order from most to least similar
             return list(most_similar_users)
```

```
In [28]: # Do a spot check of your function
         print("The 10 most similar users to user 1 are: {}".format(
             find_similar_users(1)[:10]))
         print("The 5 most similar users to user 3933 are: {}".format(
             find_similar_users(3933)[:5]))
         print("The 3 most similar users to user 46 are: {}".format(
             find_similar_users(46)[:3]))
```

```
The 10 most similar users to user 1 are: [3933, 23, 3782, 203, 4459, 3870, 131, 4201, 46, 5041]
The 5 most similar users to user 3933 are: [1, 23, 3782, 203, 4459]
The 3 most similar users to user 46 are: [4201, 3782, 23]
```

3. Now that you have a function that provides the most similar users to each user, you will want to use these users to find articles you can recommend. Complete the functions below to return the articles you would recommend to each user.

```
In [29]: def get_article_names(article_ids, df=df):
             '''
             INPUT:
             article_ids - (list) a list of article ids
             df - (pandas dataframe) df as defined at the top of the notebook

             OUTPUT:
             article_names - (list) a list of article names associated with the list of article
                             (this is identified by the title column)
             '''

             article_ids = list(np.array(article_ids).astype(float))
             article_names = list(df[df['article_id'].isin(
                 article_ids)].drop_duplicates('article_id')['title'])

             return article_names  # Return the article names associated with list of article id
```

```
In [30]: def get_user_articles(user_id, user_item=user_item):
             '''
             INPUT:
             user_id - (int) a user id
             user_item - (pandas dataframe) matrix of users by articles:
                         1's when a user has interacted with an article, 0 otherwise

             OUTPUT:
             article_ids - (list) a list of the article ids seen by the user
             article_names - (list) a list of article names associated with the list of article
                             (this is identified by the doc_full_name column in df_content)

             Description:
             Provides a list of the article_ids and article titles that have been seen by a user
             '''
```

```
    user_items = user_item.loc[user_id]
    article_ids = list(user_items[user_items > 0].index)
    article_ids = list(np.array(article_ids).astype(str))
    article_names = get_article_names(article_ids)

    return article_ids, article_names  # return the ids and names
```

In [31]:
```
def user_user_recs(user_id, m=10):
    '''
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them as rec
    Does this until m recommendations are found

    Notes:
    Users who are the same closeness are chosen arbitrarily as the 'next' user

    For the user where the number of recommended articles starts below m
    and ends exceeding m, the last items are chosen arbitrarily

    '''

    recs = []

    articles_seen_ids, articles_seen_names = get_user_articles(user_id)

    most_similar_users = find_similar_users(user_id)

    for user in most_similar_users:
        similar_article_ids, similar_article_names = get_user_articles(user)
        rec_article_ids = np.setdiff1d(
            similar_article_ids, articles_seen_ids, assume_unique=True)
        recs.extend(rec_article_ids)
        recs = list(np.unique(recs))

        if len(recs) >= m:
            break

    return recs[:m]  # return your recommendations for this user_id
```

In [32]:
```
# Check Results
```

```
# Return 10 recommendations for user 1
get_article_names(user_user_recs(1, 10))
```

Out[32]:
```
['analyze energy consumption in buildings',
 'analyze accident reports on amazon emr spark',
 '520    using notebooks with pixiedust for fast, flexi...\nName: title, dtype: object'
 '1448    i ranked every intro to data science course on...\nName: title, dtype: object
 'data tidying in data science experience',
 'airbnb data for analytics: vancouver listings',
 'recommender systems: approaches & algorithms',
 'airbnb data for analytics: mallorca reviews',
 'analyze facebook data using ibm watson and watson studio',
 'a tensorflow regression model to predict house values']
```

In [33]:
```
# Test your functions here - No need to change this code - just run this cell
assert set(get_article_names(['1024.0', '1176.0', '1305.0', '1314.0', '1422.0', '1427.0

assert set(get_article_names(['1320.0', '232.0', '844.0'])) == set(['housing (2015): un
                                                                     'use the cloudant-s
assert set(get_user_articles(20)[0]) == set(['1320.0', '232.0', '844.0'])
assert set(get_user_articles(20)[1]) == set(['housing (2015): united states demographic
                                              'self-service data preparation with ibm da
assert set(get_user_articles(2)[0]) == set(
    ['1024.0', '1176.0', '1305.0', '1314.0', '1422.0', '1427.0'])
assert set(get_user_articles(2)[1]) == set(['using deep learning to reconstruct high-re
                                             'healthcare python streaming application de
print("If this is all you see, you passed all of our tests!  Nice job!")
```

```
If this is all you see, you passed all of our tests!  Nice job!
```

4. Now we are going to improve the consistency of the **user_user_recs** function from above.

- Instead of arbitrarily choosing when we obtain users who are all the same closeness to a given user - choose the users that have the most total article interactions before choosing those with fewer article interactions.
- Instead of arbitrarily choosing articles from the user where the number of recommended articles starts below m and ends exceeding m, choose articles with the articles with the most total interactions before choosing those with fewer total interactions. This ranking should be what would be obtained from the **top_articles** function you wrote earlier.

In [34]:
```
def get_top_sorted_users(user_id, df=df, user_item=user_item):
    '''
    INPUT:
    user_id - (int)
    df - (pandas dataframe) df as defined at the top of the notebook
    user_item - (pandas dataframe) matrix of users by articles:
            1's when a user has interacted with an article, 0 otherwise
```

```
        OUTPUT:
        neighbors_df - (pandas dataframe) a dataframe with:
                        neighbor_id - is a neighbor user_id
                        similarity - measure of the similarity of each user to the provided
                        num_interactions - the number of articles viewed by the user - if a

        Other Details - sort the neighbors_df by the similarity and then by number of inter
                        highest of each is higher in the dataframe

        '''

        similarity = user_item.loc[user_id].dot(user_item.transpose())
        similarity.rename('similarity', inplace=True)

        num_interactions = df.groupby('user_id')['article_id'].count()
        num_interactions.rename('num_interactions', inplace=True)

        neighbors_df = similarity.to_frame().join(num_interactions)
        neighbors_df.index.rename('neighbor_id', inplace=True)

        neighbors_df.sort_values(
            by=['similarity', 'num_interactions'], ascending=False, inplace=True)
        neighbors_df.drop(user_id, inplace=True)

        return neighbors_df  # Return the dataframe specified in the doc_string
```

In [35]:
```
# top articles
top_articles = df.groupby('article_id')[
    'user_id'].count().sort_values(ascending=False)
top_articles
```

Out[35]:
```
article_id
1429.0    937
1330.0    927
1431.0    671
1427.0    643
1364.0    627
         ...
1113.0      1
1119.0      1
984.0       1
1127.0      1
1266.0      1
Name: user_id, Length: 714, dtype: int64
```

In [36]:
```
def user_user_recs_part2(user_id, m=10):
    '''
```

```
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user by article id
    rec_names - (list) a list of recommendations for the user by article title

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and provides them as rec
    Does this until m recommendations are found

    Notes:
    * Choose the users that have the most total article interactions
    before choosing those with fewer article interactions.

    * Choose articles with the articles with the most total interactions
    before choosing those with fewer total interactions.

    '''
    recs = []

    articles_seen_ids, articles_seen_names = get_user_articles(user_id)

    most_similar_users = get_top_sorted_users(user_id).index

    for user in most_similar_users:
        similar_article_ids, similar_article_names = get_user_articles(user)
        rec_article_ids = np.setdiff1d(
            similar_article_ids, articles_seen_ids, assume_unique=True)
        # arrange article ids by total interactions
        rec_article_ids = list(
            top_articles[rec_article_ids.astype(float)].sort_values(ascending=False).in
        recs.extend(rec_article_ids)
        recs = list(np.unique(recs))

        if len(recs) >= m:
            break

    recs = recs[:m]
    rec_names = get_article_names(recs)

    return recs, rec_names
```

In [37]:
```
# Quick spot check - don't change this code - just use it to test your functions
rec_ids, rec_names = user_user_recs_part2(20, 10)
print("The top 10 recommendations for user 20 are the following article ids:")
```

```
print(rec_ids)
print()
print("The top 10 recommendations for user 20 are the following article names:")
print(rec_names)
```

```
The top 10 recommendations for user 20 are the following article ids:
[12.0, 109.0, 125.0, 142.0, 164.0, 205.0, 302.0, 336.0, 362.0, 465.0]

The top 10 recommendations for user 20 are the following article names:
['timeseries data analysis of iot events by using jupyter notebook', 'dsx: hybrid mode', 'accele
```

5. Use your functions from above to correctly fill in the solutions to the dictionary below. Then test your dictionary against the solution. Provide the code you need to answer each following the comments below.

```
In [38]: # Tests with a dictionary of results

         # Find the user that is most similar to user 1
         user1_most_sim = get_top_sorted_users(1).index[0]
         # Find the 10th most similar user to user 131
         user131_10th_sim = get_top_sorted_users(131).index[9]
```

```
In [39]: # Dictionary Test Here
         sol_5_dict = {
             'The user that is most similar to user 1.': user1_most_sim,
             'The user that is the 10th most similar to user 131': user131_10th_sim,
         }

         t.sol_5_test(sol_5_dict)
```

```
This all looks good!  Nice job!
```

6. If we were given a new user, which of the above functions would you be able to use to make recommendations? Explain. Can you think of a better way we might make recommendations? Use the cell below to explain a better method for new users.

Since the new user id is not present in the user_item dataframe, we cannot make recommendations using collaborative filtering. In other words, we have no information about the user to find other similar users. Hence, we will use a rank based recommender system to recommend articles to the new user.

7. Using your existing functions, provide the top 10 recommended articles you would provide for the a new user below. You can test your function against our thoughts to make sure we are all on the same page with how we might make a recommendation.

```
In [40]: new_user = '0.0'

         # What would your recommendations be for this new user '0.0'?  As a new user, they have
         # Provide a list of the top 10 article ids you would give to
         # Your recommendations here
         new_user_recs = list(np.array(get_top_article_ids(10)).astype(str))
```

```
In [41]: assert set(new_user_recs) == set(['1314.0', '1429.0', '1293.0', '1427.0', '1162.0', '13
                                              ), "Oops!  It makes sense that in this case we would w

         print("That's right!  Nice job!")

That's right!  Nice job!
```

1.1.4 Part IV: Content Based Recommendations (EXTRA - NOT REQUIRED)

Another method we might use to make recommendations is to perform a ranking of the highest ranked articles associated with some term. You might consider content to be the **doc_body**, **doc_description**, or **doc_full_name**. There isn't one way to create a content based recommendation, especially considering that each of these columns hold content related information.

`1.` Use the function body below to create a content based recommender. Since there isn't one right answer for this recommendation tactic, no test functions are provided. Feel free to change the function inputs if you decide you want to try a method that requires more input values. The input values are currently set with one idea in mind that you may use to make content based recommendations. One additional idea is that you might want to choose the most popular recommendations that meet your 'content criteria', but again, there is a lot of flexibility in how you might make these recommendations.

1.1.5 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [42]: def make_content_recs():
             '''
             INPUT:

             OUTPUT:

             '''
```

`2.` Now that you have put together your content-based recommendation system, use the cell below to write a summary explaining how your content based recommender works. Do you see any possible improvements that could be made to your function? Is there anything novel about your content based recommender?

1.1.6 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

Write an explanation of your content based recommendation system here.

`3.` Use your content-recommendation system to make recommendations for the below scenarios based on the comments. Again no tests are provided here, because there isn't one right answer that could be used to find these content based recommendations.

1.1.7 This part is NOT REQUIRED to pass this project. However, you may choose to take this on as an extra way to show off your skills.

```
In [43]: # make recommendations for a brand new user


         # make a recommendations for a user who only has interacted with article id '1427.0'
```

1.1.8 Part V: Matrix Factorization

In this part of the notebook, you will build use matrix factorization to make article recommendations to the users on the IBM Watson Studio platform.

1. You should have already created a **user_item** matrix above in **question 1** of **Part III** above. This first question here will just require that you run the cells to get things set up for the rest of **Part V** of the notebook.

```
In [44]: # Load the matrix here
         user_item_matrix = pd.read_pickle('user_item_matrix.p')

In [45]: # quick look at the matrix
         user_item_matrix.head()

Out[45]: article_id  0.0  100.0  1000.0  1004.0  1006.0  1008.0  101.0  1014.0  1015.0  \
         user_id
         1           0.0    0.0     0.0     0.0     0.0     0.0    0.0     0.0     0.0
         2           0.0    0.0     0.0     0.0     0.0     0.0    0.0     0.0     0.0
         3           0.0    0.0     0.0     0.0     0.0     0.0    0.0     0.0     0.0
         4           0.0    0.0     0.0     0.0     0.0     0.0    0.0     0.0     0.0
         5           0.0    0.0     0.0     0.0     0.0     0.0    0.0     0.0     0.0

         article_id  1016.0  ...  977.0  98.0  981.0  984.0  985.0  986.0  990.0  \
         user_id             ...
         1              0.0  ...    0.0   0.0    1.0    0.0    0.0    0.0    0.0
         2              0.0  ...    0.0   0.0    0.0    0.0    0.0    0.0    0.0
         3              0.0  ...    1.0   0.0    0.0    0.0    0.0    0.0    0.0
         4              0.0  ...    0.0   0.0    0.0    0.0    0.0    0.0    0.0
         5              0.0  ...    0.0   0.0    0.0    0.0    0.0    0.0    0.0

         article_id  993.0  996.0  997.0
         user_id
         1             0.0    0.0    0.0
         2             0.0    0.0    0.0
         3             0.0    0.0    0.0
         4             0.0    0.0    0.0
         5             0.0    0.0    0.0

         [5 rows x 714 columns]
```

2. In this situation, you can use Singular Value Decomposition from numpy on the user-item matrix. Use the cell to perform SVD, and explain why this is different than in the lesson.

```
In [46]: # Perform SVD on the User-Item Matrix Here

         # use the built in to get the three matrices
         u, s, vt = np.linalg.svd(user_item_matrix)
```

Provide your response here.

`3.` Now for the tricky part, how do we choose the number of latent features to use? Running the below cell, you can see that as the number of latent features increases, we obtain a lower error rate on making predictions for the 1 and 0 values in the user-item matrix. Run the cell below to get an idea of how the accuracy improves as we increase the number of latent features.

```
In [47]: num_latent_feats = np.arange(10, 700+10, 20)
         sum_errs = []

         for k in num_latent_feats:
             # restructure with k latent features
             s_new, u_new, vt_new = np.diag(s[:k]), u[:, :k], vt[:k, :]

             # take dot product
             user_item_est = np.around(np.dot(np.dot(u_new, s_new), vt_new))

             # compute error for each prediction to actual value
             diffs = np.subtract(user_item_matrix, user_item_est)

             # total errors and keep track of them
             err = np.sum(np.sum(np.abs(diffs)))
             sum_errs.append(err)


         plt.plot(num_latent_feats, 1 - np.array(sum_errs)/df.shape[0])
         plt.xlabel('Number of Latent Features')
         plt.ylabel('Accuracy')
         plt.title('Accuracy vs. Number of Latent Features');
```

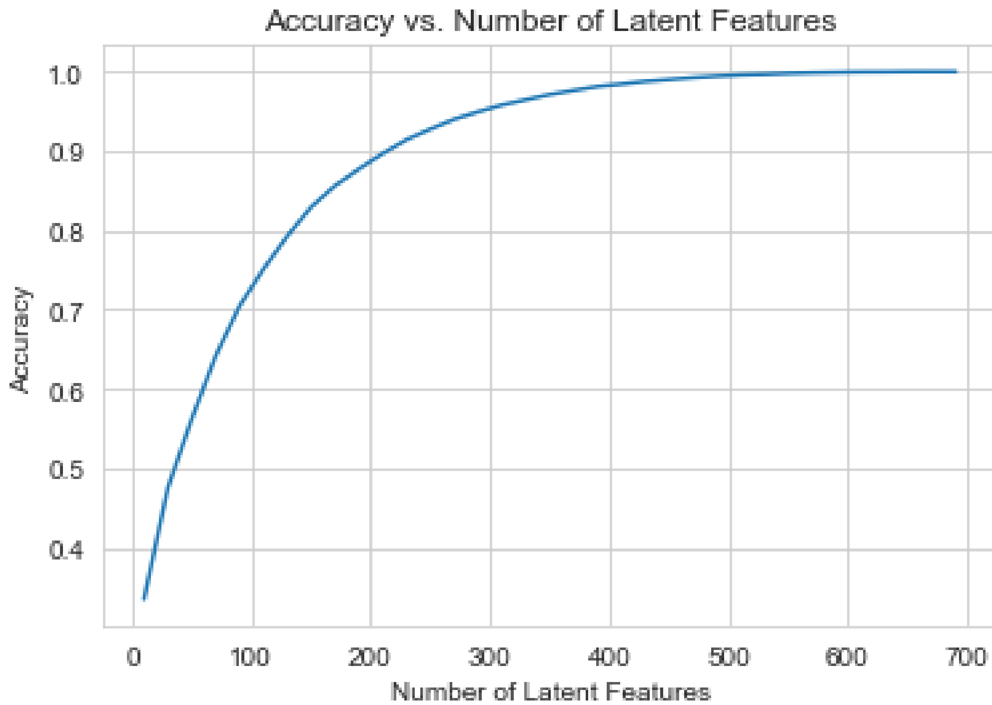

4. From the above, we can't really be sure how many features to use, because simply having a better way to predict the 1's and 0's of the matrix doesn't exactly give us an indication of if we are able to make good recommendations. Instead, we might split our dataset into a training and test set of data, as shown in the cell below.

Use the code from question 3 to understand the impact on accuracy of the training and test sets of data with different numbers of latent features. Using the split below:

- How many users can we make predictions for in the test set?
- How many users are we not able to make predictions for because of the cold start problem?
- How many articles can we make predictions for in the test set?
- How many articles are we not able to make predictions for because of the cold start problem?

```
In [48]: df_train = df.head(40000)
         df_test = df.tail(5993)


         def create_test_and_train_user_item(df_train, df_test):
             '''
             INPUT:
             df_train - training dataframe
             df_test - test dataframe

             OUTPUT:
```

```
        user_item_train - a user-item matrix of the training dataframe
                          (unique users for each row and unique articles for each column)
        user_item_test - a user-item matrix of the testing dataframe
                        (unique users for each row and unique articles for each column)
        test_idx - all of the test user ids
        test_arts - all of the test article ids

        '''

        user_item_train = create_user_item_matrix(df_train)
        user_item_test = create_user_item_matrix(df_test)

        test_idx = user_item_test.index
        test_arts = user_item_test.columns

        return user_item_train, user_item_test, test_idx, test_arts


    user_item_train, user_item_test, test_idx, test_arts = create_test_and_train_user_item(
        df_train, df_test)
```

In [49]:
```
# How many users can we make predictions for in the test set?
len(np.intersect1d(test_idx, user_item_train.index))
```

Out[49]: 20

In [50]:
```
# How many users in the test set are we not able to make predictions for because of the
len(test_idx) - len(np.intersect1d(test_idx, user_item_train.index))
```

Out[50]: 662

In [51]:
```
# How many movies can we make predictions for in the test set?
len(np.intersect1d(test_arts, user_item_train.columns))
```

Out[51]: 574

In [52]:
```
# How many movies in the test set are we not able to make predictions for because of th
len(test_arts) - len(np.intersect1d(test_arts, user_item_train.columns))
```

Out[52]: 0

In [53]:
```
# Replace the values in the dictionary below
a = 662
b = 574
c = 20
d = 0


sol_4_dict = {
```

```
    'How many users can we make predictions for in the test set?': c,
    'How many users in the test set are we not able to make predictions for because of
    'How many movies can we make predictions for in the test set?':  b,
    'How many movies in the test set are we not able to make predictions for because of
}

t.sol_4_test(sol_4_dict)
```

```
Awesome job!  That's right!  All of the test movies are in the training data, but there are only
```

5. Now use the **user_item_train** dataset from above to find U, S, and V transpose using SVD. Then find the subset of rows in the **user_item_test** dataset that you can predict using this matrix decomposition with different numbers of latent features to see how many features makes sense to keep based on the accuracy on the test data. This will require combining what was done in questions 2 - 4.

Use the cells below to explore how well SVD works towards making predictions for recommendations on the test data.

In [54]:
```
# fit SVD on the user_item_train matrix
# fit svd similar to above then use the cells below
u_train, s_train, vt_train = np.linalg.svd(user_item_train)
```

In [55]:
```
# Use these cells to see how well you can use the training
# decomposition to predict on test data
```

In [56]:
```
# Users can we make predictions for in the test set
pred_users = np.intersect1d(test_idx, user_item_train.index)
```

In [57]:
```
# Test user indices
idx = np.where(user_item_train.index.isin(pred_users))[0]
```

In [58]:
```
# Test item indices
col = np.where(user_item_train.columns.isin(user_item_test.columns))[0]
```

In [59]:
```
num_latent_feats = np.arange(10, 700+10, 20)
sum_errs = []

for k in num_latent_feats:
    # restructure with k latent features
    s_new, u_new, vt_new = np.diag(
        s_train[:k]), u_train[:, :k], vt_train[:k, :]

    # take dot product
    user_item_est = np.around(
        np.dot(np.dot(u_new[idx, :], s_new), vt_new[:, col]))

    # compute error for each prediction to actual value
    diffs = np.subtract(user_item_test.loc[pred_users], user_item_est)
```

```
        # total errors and keep track of them
        err = np.sum(np.sum(np.abs(diffs)))
        sum_errs.append(err)


    plt.plot(num_latent_feats, 1 - np.array(sum_errs)/(len(idx) * len(col)))
    plt.xlabel('Number of Latent Features')
    plt.ylabel('Accuracy')
    plt.title('Accuracy vs. Number of Latent Features')
```

Out[59]: Text(0.5, 1.0, 'Accuracy vs. Number of Latent Features')

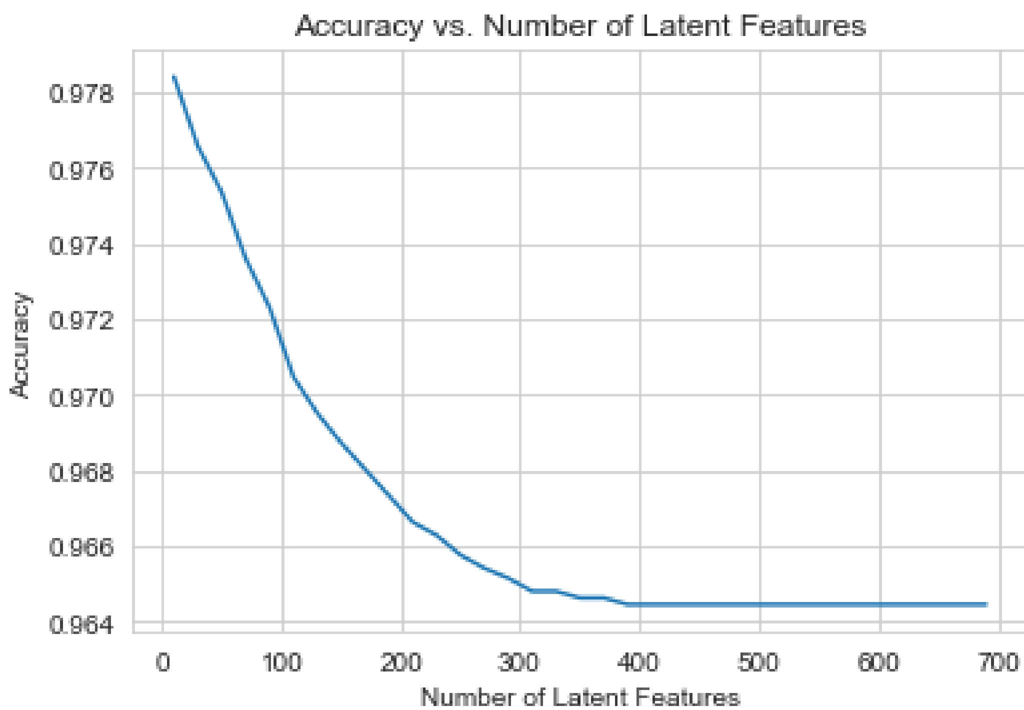

6. Use the cell below to comment on the results you found in the previous question. Given the circumstances of your results, discuss what you might do to determine if the recommendations you make with any of the above recommendation systems are an improvement to how users currently find articles?

From the plot above, we can see how the accuracy goes down by adding more latent features. There are 4487 user ids in the training set and 682 user ids in the test set. However, they only have 20 user ids in common. The observation in the training data are not representative of the population as a whole. Increasing the order of randomness before splitting the data may help improve this. Secondly, the type of ratings in the user-item matrix is binary. Have a rating system with a broader scale may help improve predictions.

We could do an A/B test online for a period of time to measure the effectiveness of any of the above recommendation systems vs. the existing system in place. We can create a cookie-

based diversion to split users into 2 experimental groups and use the number of articles read by users as the performance evaluation metric.

Extras Using your workbook, you could now save your recommendations for each user, develop a class to make new predictions and update your results, and make a flask app to deploy your results. These tasks are beyond what is required for this project. However, from what you learned in the lessons, you certainly capable of taking these tasks on to improve upon your work here!

1.2 Conclusion

Congratulations! You have reached the end of the Recommendations with IBM project!

Tip: Once you are satisfied with your work here, check over your report to make sure that it is satisfies all the areas of the rubric. You should also probably remove all of the "Tips" like this one so that the presentation is as polished as possible.

1.3 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File** > **Download as** submenu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

```
In [ ]: from subprocess import call
        call(['python', '-m', 'nbconvert', 'Recommendations_with_IBM.ipynb'])
```

```
In [ ]:
```